

Tidy Data Exploration

Tai Chou-Kudu

Overview

Part 1: Airplane flight delays

Consider the following dataset:

		Los_Angeles	Phoenix	San_Diego	San_Francisco	Seattle
ALASKA	On_Time	497	221	212	503	1841
	Delayed	62	12	20	102	305
AM	On_Time	694	4840	383	320	301
WEST	Delayed	117	415	65	129	61

The above table describes arrival delays for two different airlines across several destinations. The numbers correspond to the number of flights that were in either the delayed category or the on time category.

Problem 1: Read the information from `flightdelays.csv` into R, and use `tidyr` and `dplyr` to convert this data into a tidy/tall format with names and complete data for all columns. Your final data frame should have `City`, `On_Time_Flights` and `Delayed_Flights` as columns (the exact names are up to you). In addition to `pivot_longer`, `pivot_wider` and `rename`, you might find the `tidyr` function `fill` helpful for completing this task efficiently. Although this is a small dataset that you could easily reshape by hand, you should solve this problem using tidyverse functions that do the work for you.

```
library(tidyr)
library(dplyr)
library(ggplot2)
library(readr)
library(stringr)
library(here)
```

Read Flight Delays Data

Clicked “raw” data on github to retrieve URL for dataset, not just URL for gitHub HTML page. Reading data and inspecting data’s column info below.

```
flightdelays <- readr::read_csv("https://raw.githubusercontent.com/georgehagstrom/DATA607/master/flightdelays.csv")
spec(flightdelays)
```

```
cols(
  ...1 = col_character(),
  ...2 = col_character(),
  Los_Angeles = col_double(),
  Phoenix = col_double(),
  San_Diego = col_double(),
  San_Francisco = col_double(),
  Seattle = col_double()
)
```

Tidy Flight Delays Dataset

Renaming columns to give Airlines a column name, filling in missing airline data, performing a tall pivot.

```
tall_flightdelays <- flightdelays %>%
  rename(Airline = ...1,
         Flight_Status = ...2) %>%
  fill(Airline, .direction = "down") %>%
  pivot_longer(
    cols = Los_Angeles:Seattle,
    names_to = "City",
    values_to = "Flights"
  )

tall_flightdelays <- tall_flightdelays %>%
  mutate(Airline = Airline %>%
    str_to_lower() %>%
    str_to_title() %>%
    str_replace_all(" ", "_")
  )
```

```

final_flightdelays <- tall_flightdelays %>%
  pivot_wider(
    names_from = Flight_Status,
    values_from = Flights
  ) %>%
  rename(
    On_Time_Flights = On_Time,
    Delayed_Flights = Delayed
  ) %>%
  arrange(City) %>%
  relocate(City, .before = Airline)

print(final_flightdelays)

```

```

# A tibble: 10 x 4
  City      Airline On_Time_Flights Delayed_Flights
  <chr>    <chr>      <dbl>         <dbl>
1 Los_Angeles Alaska        497           62
2 Los_Angeles Am_West        694          117
3 Phoenix   Alaska        221           12
4 Phoenix   Am_West       4840          415
5 San_Diego  Alaska        212           20
6 San_Diego  Am_West       383           65
7 San_Francisco Alaska        503          102
8 San_Francisco Am_West       320          129
9 Seattle    Alaska       1841          305
10 Seattle    Am_West       301           61

```

Problem 2: Take the data-frame that you tidied and cleaned in Problem 1 and create additional columns which contain the fraction of on-time and delayed flights at each airport. Then create a Cleveland Multiway Dot Plot (see [this tutorial page for a description for how](#)) to visualize the difference in flight delays between the two airlines at each city in the dataset. Compare the airlines and airports using the dot-plot- what are your conclusions?

Add Columns and Visualize Flight Delays

Add calculated percentage columns, visualize flight delays with a Cleveland Multiway Dot Plot.

```

final_flightdelays <- final_flightdelays %>%
  group_by(City) %>%
  mutate(
    Total_Flights = On_Time_Flights + Delayed_Flights,
    On_Time_Pct= round(On_Time_Flights / Total_Flights, 2),
    Delayed_Pct= round(Delayed_Flights / Total_Flights, 2)
  ) %>%
  ungroup()

final_flightdelays

```

```

# A tibble: 10 x 7
  City      Airline On_Time_Flights Delayed_Flights Total_Flights On_Time_Pct
  <chr>    <chr>         <dbl>          <dbl>         <dbl>         <dbl>
1 Los_Angeles Alaska         497             62           559           0.89
2 Los_Angeles Am_West         694            117           811           0.86
3 Phoenix   Alaska         221             12           233           0.95
4 Phoenix   Am_West        4840            415          5255           0.92
5 San_Diego  Alaska         212             20           232           0.91
6 San_Diego  Am_West         383             65           448           0.85
7 San_Franci~ Alaska         503            102           605           0.83
8 San_Franci~ Am_West         320            129           449           0.71
9 Seattle    Alaska        1841            305          2146           0.86
10 Seattle   Am_West         301             61           362           0.83
# i 1 more variable: Delayed_Pct <dbl>

```

```


final_flightdelays <- final_flightdelays %>%
  mutate(City = factor(City, levels = unique(City)),
    Airline = factor(Airline, levels = unique(Airline)))

# Create the Cleveland Multiway Dot Plot
ggplot(final_flightdelays, aes(x = On_Time_Pct, y = City, color = Airline)) +
  geom_point(size = 3) +
  geom_line(aes(group = City), linetype = "solid") +
  labs(
    title = "Comparison of On-Time Flight Percentages by Airline and City",
    x = "On-Time Flight Percentage",
    y = "City",
    color = "Airline"
  ) +
  theme_minimal() +

```

```
theme(
  axis.title = element_text(size = 12, margin = margin(t = 30)),
  axis.text = element_text(size = 10),
  legend.title = element_text(size = 12),
  legend.text = element_text(size = 10),
  plot.title = element_text(size = 12, face = "bold", hjust = 0.5)
)
```

Comparison of On-Time Flight Percentages by Airline and City



Optional: If you want to make a fancier visualization consider adding text labels containing the airline names above the dots using `geom_text` and `position = position_nudge(...)` with appropriate arguments.

Part 2: Mixed Drink Recipes

In the second part of this assignment we will be working with a dataset containing ingredients for different types of mixed drinks. This dataset is untidy and messy- it is in a wide data format and contains some inconsistencies that should be fixed.

Problem 3: Load the mixed drink recipe dataset into R from the file `MixedDrinkRecipes-prep.csv`, which you can download from my github page by [clicking here](#). The variables `ingredient1` through `ingredient6` list the ingredients of the cocktail listed in the `name` column. Notice that there are many NA values in the ingredient columns, indicating that most cocktails have under 6 ingredients.

Tidy this dataset using `pivot_longer` to create a new data frame where each there is a row corresponding to each ingredient of all the cocktails, and an additional variable specifying the “rank” of that cocktail in the original recipe, i.e. it should look like this:

name	category	Ingredient_Rank	Ingredient
Gauguin	Cocktail Classics	1	Light Rum
Gauguin	Cocktail Classics	2	Passion Fruit Syrup
Gauguin	Cocktail Classics	3	Lemon Juice
Gauguin	Cocktail Classics	4	Lime Juice
Fort Lauderdale	Cocktail Classics	1	Light Rum

where the data-type of `Ingredient_Rank` is an integer. Hint: Use the `parse_number()` function in `mutate` after your initial pivot.

Read Mixed Drinks Data

Read data using `readr` package, use `here::here` to specify file local to the project root instead of copying a full pathname/ device-specific local path.

```
mixed_drinks <- readr::read_csv(here::here("MixedDrinkRecipes-Prep.csv"))
```

Tidy Mixed Drinks Dataset

Pivot longer to display tidy dataset where each ingredient of all the drinks has a row, including the rank in the original recipe and dataset.

```
tidy_drinks <- mixed_drinks %>%
  pivot_longer(
    cols = ingredient1:ingredient6,
    names_to = "Ingredient_Rank",
    values_to = "Ingredient",
    names_prefix = "ingredient") %>%
    mutate(Ingredient_Rank =
      parse_number(Ingredient_Rank)) %>%
      filter(!is.na(Ingredient))
  )

print(tidy_drinks)
```

```
# A tibble: 3,934 x 4
  name          category      Ingredient_Rank Ingredient
  <chr>         <chr>          <dbl> <chr>
1 Gauguin      Cocktail Classics      1 Light Rum
2 Gauguin      Cocktail Classics      2 Passion Fruit Syrup
3 Gauguin      Cocktail Classics      3 Lemon Juice
4 Gauguin      Cocktail Classics      4 Lime Juice
5 Fort Lauderdale Cocktail Classics      1 Light Rum
6 Fort Lauderdale Cocktail Classics      2 Sweet Vermouth
7 Fort Lauderdale Cocktail Classics      3 Juice of Orange
8 Fort Lauderdale Cocktail Classics      4 Juice of a Lime
9 Apple Pie    Cordials and Liqueurs      1 Apple schnapps
10 Apple Pie   Cordials and Liqueurs      2 Cinnamon schnapps
# i 3,924 more rows
```

Problem 4: Some of the ingredients in the ingredient list have different names, but are nearly the same thing. An example of such a pair is **Lemon Juice** and **Juice of a lemon**, which are considered different ingredients in this dataset, but which perhaps should be treated as the same depending on the analysis you are doing. Make a list of the ingredients appearing in the ingredient list ranked by how commonly they occur along with the number of occurrences, and print the first 10 elements of the list here. Then check more ingredients (I suggest looking at more ingredients and even sorting them alphabetically using `arrange(asc(ingredient))`) and see if you can spot pairs of ingredients that are similar but have different names. Use `if_else` ([click here for if_else](#)) or `case_when` in combination with `mutate` to make it so that the pairs of ingredients you found have the same name. You don't have to find all pairs, but find at least 5 pairs of ingredients to rename. Because the purpose of this renaming is to facilitate a hypothetical future analysis, you can choose your own criteria for similarity as long as it is somewhat justifiable.

Notice that there are some ingredients that appear to be two or more ingredients strung together with commas. These would be candidates for more cleaning though this exercise doesn't ask you to fix them.

Standardize Data and Count Top 10 Ingredients

Use **Case When** to standardize at least 5 pairs of ingredients, sort by the top 10 from the standardized dataset.

```
ingredient_counts <- tidy_drinks %>%
  mutate(
    Ingredient = case_when(
      Ingredient == "Fresh lemon juice" ~ "Lemon Juice",
```

```

    Ingredient == "Juice of a Lime" ~ "Fresh Lime Juice",
    Ingredient == "Juice of Orange" ~ "Fresh orange juice",
    Ingredient == "ginger ale" ~ "Ginger ale",
    Ingredient == "Juice of a Lemon" ~ "Lemon Juice",
    TRUE ~ Ingredient
  )
) %>%
  count(Ingredient, sort = TRUE)

print(head(ingredient_counts, 10))

```

```

# A tibble: 10 x 2
  Ingredient      n
  <chr>         <int>
1 Lemon Juice   268
2 Gin           176
3 Fresh Lime Juice 142
4 Simple Syrup  115
5 Light Rum     114
6 Vodka         114
7 Dry Vermouth  107
8 Triple Sec    107
9 Powdered Sugar  90
10 Grenadine     85

```

Problem 5: Some operations are easier to do on **wide** data rather than **tall** data. Find the 10 most common pairs of ingredients occurring in the top 2 ingredients in a recipe. It is much easier to do this with a **wide** dataset, so use `pivot_wider` to change the data so that each row contains all of the ingredients of a single cocktail, just like in the format of the original data-set. Then use `count` on the 1 and 2 columns to determine the most common pairs (see chapter 3 for a refresher on `count`).

Note: You may be interested to read about the `widyr` package here: [widyr page](#). It is designed to solve problems like this one and uses internal pivot steps to accomplish it so that the final result is tidy. I'm actually unaware of any easy ways of solving problem 5 without pivoting to a wide dataset.

Count Top 10 Ingredient Pairs

Pivot wider so that we can see each drink and their 1 and 2 ranked ingredients. Count and sort to see the top 10 ingredient pairs.


```

wide_drinks <- tidy_drinks %>%
  filter(Ingredient_Rank <= 2) %>%
  pivot_wider(
    names_from = Ingredient_Rank,
    values_from = Ingredient,
    names_prefix = "Ingredient_"
  )

ingredient_pairs <- wide_drinks %>%
  count(Ingredient_1, Ingredient_2, sort = TRUE)

print(head(ingredient_pairs, 10))

```

```

# A tibble: 10 x 3
  Ingredient_1      Ingredient_2      n
  <chr>           <chr>          <int>
1 Gin             Dry Vermouth    23
2 Juice of a Lemon Powdered Sugar  23
3 Whole Egg       Powdered Sugar  13
4 Light Rum       Fresh Lime Juice 12
5 Gin             Triple Sec      9
6 Bourbon whiskey Fresh lemon juice 8
7 Brandy          Sweet Vermouth  7
8 Gin             Sweet Vermouth  7
9 Light Rum       Pineapple Juice 7
10 Light Rum      Sweet Vermouth  7

```